


INFERENCE per query: one LLM call, then one forward pass that scores the whole corpus

--> used during training only



● semantic channel ● graph channel ● learned in training ● frozen, no parameters ● the correct paper